

Time series prediction using RNNs, with TensorFlow and Cloud ML Engine

This notebook illustrates:

1. Creating a Recurrent Neural Network in TensorFlow
2. Creating a Custom Estimator in tf.estimator
3. Training on Cloud ML Engine

Simulate some time-series data

Essentially a set of sinusoids with random amplitudes and frequencies.

In [23]:

```
import os
PROJECT = 'cloud-training-demos' # REPLACE WITH YOUR PROJECT ID
BUCKET = 'cloud-training-demos-ml' # REPLACE WITH YOUR BUCKET NAME
REGION = 'us-centrall' # REPLACE WITH YOUR BUCKET REGION e.g. us-centrall
os.environ['TFVERSION'] = '1.8' # Tensorflow version
```

In [24]:

```
# for bash
os.environ['PROJECT'] = PROJECT
os.environ['BUCKET'] = BUCKET
os.environ['REGION'] = REGION
```

In [25]:

```
%%bash
gcloud config set project $PROJECT
gcloud config set compute/region $REGION
```

Updated property [core/project].
Updated property [compute/region].

In [26]:

```
import tensorflow as tf
print(tf.__version__)
```

1.8.0

In [27]:

```


import numpy as np
import seaborn as sns
import pandas as pd

SEQ_LEN = 10
def create_time_series():
    freq = (np.random.random() * 0.5) + 0.1 # 0.1 to 0.6
    ampl = np.random.random() + 0.5 # 0.5 to 1.5
    x = np.sin(np.arange(0, SEQ_LEN) * freq) * ampl
    return x

for i in range(0, 5):
    sns.tsplot( create_time_series() ); # 5 series

```

/usr/local/envs/py3env/lib/python3.5/site-packages/matplotlib/font_manager.py:1320: UserWarning: findfont: Font family ['sans-serif'] not found. Falling back to DejaVu Sans
(prop.get_family(), self.defaultFamily[fontext]))



In [28]:

```

def to_csv(filename, N):
    with open(filename, 'w') as ofp:
        for lineno in range(0, N):
            seq = create_time_series()
            line = ",".join(map(str, seq))
            ofp.write(line + '\n')

to_csv('train.csv', 1000) # 1000 sequences
to_csv('valid.csv', 50)

```

In [29]:

```
!head -5 train.csv valid.csv
```

```
==> train.csv <==
0.0,0.41494536271196236,0.734451551569238,0.8850308030101335,0.83204
92576193983,0.5876928930606661,0.20816469810142113,-0.21924225053918
68,-0.5962225619744843,-0.836069183382574
0.0,0.16050934484878304,0.3187737388366111,0.4725796299246527,0.6197
758245612592,0.7583035751751281,0.8862253746803034,1.001752055261012
6,1.1032678124214037,1.1893528043033446
0.0,0.2019768145665784,0.38954713901631743,0.5493320613627289,0.6699
345313575295,0.7427522829645161,0.7625914120164823,0.728036844128641
7,0.6415532683388727,0.509309337136115
0.0,0.20513560202370504,0.38685799767444984,0.5244262531017174,0.602
1389789663988,0.6111264126796244,0.5503627711095774,0.42678332861145
07,0.2544928576485325,0.05315577684418399
0.0,0.20112443738388527,0.3910248758298022,0.559103685595879,0.69598
10200217685,0.794018269517251,0.8477443436542725,0.8541609921610173,
0.8129101260008574,0.726293800975636

==> valid.csv <==
0.0,0.4518834761927216,0.8463432780938488,1.1332529187241658,1.27615
29862603307,1.256884270714006,1.077895368911241,0.7619315258995626,
0.3491442536614218,-0.10801097394423191
0.0,0.4790437748864075,0.8201581888176755,0.9251273884779285,0.76372
79675979914,0.38243108423806044,-0.1089778135632325,-0.5690091192298
001,-0.8652076640769141,-0.9122900603467594
0.0,0.3825159065598339,0.6359357547959016,0.6747323732958791,0.48581
2213690283,0.13293431847834153,-0.2648078457410006,-0.57317948049446
85,-0.6881076260362654,-0.5708049526750818
0.0,0.09017440871015896,0.1785854105056497,0.26350408280046406,0.343
2697973071416,0.4163226944765429,0.481234187353148,0.536734898327307
5,0.5817394824668976,0.6153678519953671
0.0,0.7084840320523366,1.200134233070997,1.3244794677975964,1.043463
4746979885,0.44309210797471377,-0.29288903247151693,-0.9392305365880
832,-1.2981174951417012,-1.2597113929215498
```

RNN

For more info, see:

1. <http://colah.github.io/posts/2015-08-Understanding-LSTMs/> for the theory
2. <https://www.tensorflow.org/tutorials/recurrent> for explanations
3. <https://github.com/tensorflow/models/tree/master/tutorials/rnn/ptb> for sample code

Here, we are trying to predict from 9 values of a timeseries, the tenth value.

Imports

Several tensorflow packages and shutil

In [30]:

```
import tensorflow as tf
import shutil
import tensorflow.contrib.metrics as metrics
import tensorflow.contrib.rnn as rnn
```

Input Fn to read CSV

Our CSV file structure is quite simple -- a bunch of floating point numbers (note the type of DEFAULTS). We ask for the data to be read BATCH_SIZE sequences at a time. The Estimator API in tf.contrib.learn wants the features returned as a dict. We'll just call this timeseries column 'rawdata'.

Our CSV file sequences consist of 10 numbers. We'll assume that 9 of them are inputs and we need to predict the last one.

In [31]:

```
DEFAULTS = [[0.0] for x in range(0, SEQ_LEN)]
BATCH_SIZE = 20
TIMESERIES_COL = 'rawdata'
# In each sequence, column index 0 to N_INPUTS - 1 are features, and column index
x N_INPUTS to SEQ_LEN are labels
N_OUTPUTS = 1
N_INPUTS = SEQ_LEN - N_OUTPUTS
```

Reading data using the Estimator API in tf.estimator requires an input_fn. This input_fn needs to return a dict of features and the corresponding labels.

So, we read the CSV file. The Tensor format here will be a scalar -- entire line. We then decode the CSV. At this point, all_data will contain a list of scalar Tensors. There will be SEQ_LEN of these tensors.

We split this list of SEQ_LEN tensors into a list of N_INPUTS Tensors and a list of N_OUTPUTS Tensors. We stack them along the first dimension to then get a vector Tensor for each. We then put the inputs into a dict and call it features. The other is the ground truth, so labels.

In [32]:

```
# Read data and convert to needed format
def read_dataset(filename, mode, batch_size = 512):
    def _input_fn():
        # Provide the ability to decode a CSV
        def decode_csv(line):
            # all_data is a list of scalar tensors
            all_data = tf.decode_csv(line, record_defaults = DEFAULTS)
            inputs = all_data[:len(all_data) - N_OUTPUTS] # first N_INPUTS values
            labels = all_data[len(all_data) - N_OUTPUTS:] # last N_OUTPUTS values

            # Convert each list of rank R tensors to one rank R+1 tensor
            inputs = tf.stack(inputs, axis = 0)
            labels = tf.stack(labels, axis = 0)

            # Convert input R+1 tensor into a feature dictionary of one R+1 tensor
            features = {TIMESERIES_COL: inputs}

            return features, labels

        # Create list of files that match pattern
        file_list = tf.gfile.Glob(filename)

        # Create dataset from file list
        dataset = tf.data.TextLineDataset(file_list).map(decode_csv)

        if mode == tf.estimator.ModeKeys.TRAIN:
            num_epochs = None # indefinitely
            dataset = dataset.shuffle(buffer_size = 10 * batch_size)
        else:
            num_epochs = 1 # end-of-input after this

        dataset = dataset.repeat(num_epochs).batch(batch_size)

        iterator = dataset.make_one_shot_iterator()
        batch_features, batch_labels = iterator.get_next()
        return batch_features, batch_labels
    return _input_fn
```

Define RNN

A recursive neural network consists of possibly stacked LSTM cells.

The RNN has one output per input, so it will have 8 output cells. We use only the last output cell, but rather use it directly, we do a matrix multiplication of that cell by a set of weights to get the actual predictions. This allows for a degree of scaling between inputs and predictions if necessary (we don't really need it in this problem).

Finally, to supply a model function to the Estimator API, you need to return a EstimatorSpec. The rest of the function creates the necessary objects.

In [33]:

```

LSTM_SIZE = 3  # number of hidden layers in each of the LSTM cells

# Create the inference model
def simple_rnn(features, labels, mode):
    # 0. Reformat input shape to become a sequence
    x = tf.split(features[TIMESERIES_COL], N_INPUTS, 1)

    # 1. Configure the RNN
    lstm_cell = rnn.BasicLSTMCell(LSTM_SIZE, forget_bias = 1.0)
    outputs, _ = rnn.static_rnn(lstm_cell, x, dtype = tf.float32)

    # Slice to keep only the last cell of the RNN
    outputs = outputs[-1]

    # Output is result of linear activation of last layer of RNN
    weight = tf.get_variable("weight", initializer=tf.initializers.random_normal,
                             shape=[LSTM_SIZE, N_OUTPUTS])
    bias = tf.get_variable("bias", initializer=tf.initializers.random_normal, shape=[N_OUTPUTS])
    predictions = tf.matmul(outputs, weight) + bias

    # 2. Loss function, training/eval ops
    if mode == tf.estimator.ModeKeys.TRAIN or mode == tf.estimator.ModeKeys.EVAL:
        loss = tf.losses.mean_squared_error(labels, predictions)
        train_op = tf.contrib.layers.optimize_loss(
            loss = loss,
            global_step = tf.train.get_global_step(),
            learning_rate = 0.01,
            optimizer = "SGD")
        eval_metric_ops = {
            "rmse": tf.metrics.root_mean_squared_error(labels, predictions)
        }
    else:
        loss = None
        train_op = None
        eval_metric_ops = None

    # 3. Create predictions
    predictions_dict = {"predicted": predictions}

    # 4. Create export outputs
    export_outputs = {"predict_export_outputs": tf.estimator.export.PredictOutput(
        outputs = predictions)}

    # 5. Return EstimatorSpec
    return tf.estimator.EstimatorSpec(
        mode = mode,
        predictions = predictions_dict,
        loss = loss,
        train_op = train_op,
        eval_metric_ops = eval_metric_ops,
        export_outputs = export_outputs)

```

Estimator

Distributed training is launched off using an Estimator. The key line here is that we use `tf.estimator.Estimator` rather than, say `tf.estimator.DNNRegressor`. This allows us to provide a `model_fn`, which will be our RNN defined above. Note also that we specify a `serving_input_fn` -- this is how we parse the input data provided to us at prediction time.

In [34]:

```
# Create functions to read in respective datasets
def get_train():
    return read_dataset(filename = 'train.csv', mode = tf.estimator.ModeKeys.TRAIN
, batch_size = 512)

def get_valid():
    return read_dataset(filename = 'valid.csv', mode = tf.estimator.ModeKeys.EVAL,
batch_size = 512)
```

In [35]:

```
# Create serving input function
def serving_input_fn():
    feature_placeholders = {
        TIMESERIES_COL: tf.placeholder(tf.float32, [None, N_INPUTS])
    }

    features = {
        key: tf.expand_dims(tensor, -1)
        for key, tensor in feature_placeholders.items()
    }
    features[TIMESERIES_COL] = tf.squeeze(features[TIMESERIES_COL], axis = [2])

    return tf.estimator.export.ServingInputReceiver(features, feature_placeholders
)
```

In [36]:

```
# Create custom estimator's train and evaluate function
def train_and_evaluate(output_dir):
    estimator = tf.estimator.Estimator(model_fn = simple_rnn,
                                      model_dir = output_dir)
    train_spec = tf.estimator.TrainSpec(input_fn = get_train(),
                                       max_steps = 1000)
    exporter = tf.estimator.LatestExporter('exporter', serving_input_fn)
    eval_spec = tf.estimator.EvalSpec(input_fn = get_valid(),
                                     steps = None,
                                     exporters = exporter)
    tf.estimator.train_and_evaluate(estimator, train_spec, eval_spec)
```

In [37]:

```
# Run the model  
shutil.rmtree('outputdir', ignore_errors = True) # start fresh each time  
train_and_evaluate('outputdir')
```



```

INFO:tensorflow:Using default config.
INFO:tensorflow:Using config: {'_model_dir': 'outputdir', '_service': None, '_keep_checkpoint_every_n_hours': 10000, '_task_type': 'worker', '_num_worker_replicas': 1, '_keep_checkpoint_max': 5, '_num_ps_replicas': 0, '_global_id_in_cluster': 0, '_log_step_count_steps': 100, '_save_checkpoints_secs': 600, '_evaluation_master': '', '_train_distribute': None, '_save_summary_steps': 100, '_tf_random_seed': None, '_cluster_spec': <tensorflow.python.training.server_lib.ClusterSpec object at 0x7f65a2bc30b8>, '_master': '', '_task_id': 0, '_session_config': None, '_save_checkpoints_steps': None, '_is_chief': True}
INFO:tensorflow:Running training and evaluation locally (non-distributed).
INFO:tensorflow:Start train and evaluate loop. The evaluate will happen after 600 secs (eval_spec.throttle_secs) or training is finished.
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Saving checkpoints for 1 into outputdir/model.ckpt.
INFO:tensorflow:step = 1, loss = 2.1511927
INFO:tensorflow:global_step/sec: 13.8533
INFO:tensorflow:step = 101, loss = 0.5264278 (7.220 sec)
INFO:tensorflow:global_step/sec: 14.175
INFO:tensorflow:step = 201, loss = 0.42215365 (7.055 sec)
INFO:tensorflow:global_step/sec: 14.2523
INFO:tensorflow:step = 301, loss = 0.34791386 (7.017 sec)
INFO:tensorflow:global_step/sec: 14.7247
INFO:tensorflow:step = 401, loss = 0.26609486 (6.791 sec)
INFO:tensorflow:global_step/sec: 14.6274
INFO:tensorflow:step = 501, loss = 0.21945082 (6.836 sec)
INFO:tensorflow:global_step/sec: 14.4637
INFO:tensorflow:step = 601, loss = 0.1646782 (6.914 sec)
INFO:tensorflow:global_step/sec: 14.5217
INFO:tensorflow:step = 701, loss = 0.13758004 (6.887 sec)
INFO:tensorflow:global_step/sec: 14.123
INFO:tensorflow:step = 801, loss = 0.1219064 (7.081 sec)
INFO:tensorflow:global_step/sec: 15.7489
INFO:tensorflow:step = 901, loss = 0.10583858 (6.350 sec)
INFO:tensorflow:Saving checkpoints for 1000 into outputdir/model.ckpt.
INFO:tensorflow:Loss for final step: 0.085023135.
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Starting evaluation at 2018-09-12-20:01:59
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from outputdir/model.ckpt-1000
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Finished evaluation at 2018-09-12-20:01:59
INFO:tensorflow:Saving dict for global step 1000: global_step = 1000, loss = 0.069563285, rmse = 0.26374853
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Signatures INCLUDED in export for Classify: None
INFO:tensorflow:Signatures INCLUDED in export for Regress: None
INFO:tensorflow:Signatures INCLUDED in export for Predict: ['serving_default', 'predict_export_outputs']

```

```
INFO:tensorflow:Restoring parameters from outputdir/model.ckpt-1000
INFO:tensorflow:Assets added to graph.
INFO:tensorflow:No assets to write.
INFO:tensorflow:SavedModel written to: b"outputdir/export/exporter/t
emp-b'1536782520'/saved_model.pb"
```

Standalone Python module

To train this on Cloud ML Engine, we take the code in this notebook and make a standalone Python module.

In [38]:

```
%%bash
# Run module as-is
echo $PWD
rm -rf outputdir
export PYTHONPATH=${PYTHONPATH}:${PWD}/simplernn
python -m trainer.task \
  --train_data_paths="${PWD}/train.csv*" \
  --eval_data_paths="${PWD}/valid.csv*" \
  --output_dir=outputdir \
  --job-dir=./tmp
```

/content/datalab/training-data-analyst/courses/machine_learning/deep
dive/05_artandscience

```

/usr/local/envs/py3env/lib/python3.5/site-packages/h5py/__init__.py:
36: FutureWarning: Conversion of the second argument of issubdtype f
rom `float` to `np.floating` is deprecated. In future, it will be tr
eated as `np.float64 == np.dtype(float).type`.
    from ._conv import register_converters as _register_converters
INFO:tensorflow:Using default config.
INFO:tensorflow:Using config: {'_model_dir': 'outputdir/', '_task_ty
pe': 'worker', '_service': None, '_num_ps_replicas': 0, '_tf_random_
seed': None, '_task_id': 0, '_num_worker_replicas': 1, '_master':
'', '_is_chief': True, '_keep_checkpoint_max': 5, '_evaluation_maste
r': '', '_log_step_count_steps': 100, '_global_id_in_cluster': 0, '_
train_distribute': None, '_save_checkpoints_steps': None, '_cluster_
spec': <tensorflow.python.training.server_lib.ClusterSpec object at
0x7fa28bd35780>, '_session_config': None, '_keep_checkpoint_every_n_
hours': 10000, '_save_checkpoints_secs': 600, '_save_summary_steps':
100}
INFO:tensorflow:Running training and evaluation locally (non-distrib
uted).
INFO:tensorflow:Start train and evaluate loop. The evaluate will hap
pen after 600 secs (eval_spec.throttle_secs) or training is finishe
d.
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Graph was finalized.
2018-09-12 20:02:07.143754: I tensorflow/core/platform/cpu_feature_g
uard.cc:140] Your CPU supports instructions that this TensorFlow bin
ary was not compiled to use: AVX2 FMA
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Saving checkpoints for 1 into outputdir/model.ckpt.
INFO:tensorflow:step = 1, loss = 0.86221
INFO:tensorflow:global_step/sec: 15.0388
INFO:tensorflow:step = 101, loss = 0.30457026 (6.650 sec)
INFO:tensorflow:global_step/sec: 16.4018
INFO:tensorflow:step = 201, loss = 0.19842917 (6.097 sec)
INFO:tensorflow:global_step/sec: 19.017
INFO:tensorflow:step = 301, loss = 0.15593535 (5.258 sec)
INFO:tensorflow:global_step/sec: 19.325
INFO:tensorflow:step = 401, loss = 0.1427422 (5.175 sec)
INFO:tensorflow:global_step/sec: 19.4141
INFO:tensorflow:step = 501, loss = 0.11589954 (5.151 sec)
INFO:tensorflow:global_step/sec: 19.1968
INFO:tensorflow:step = 601, loss = 0.11280591 (5.209 sec)
INFO:tensorflow:global_step/sec: 19.0675
INFO:tensorflow:step = 701, loss = 0.099279635 (5.244 sec)
INFO:tensorflow:global_step/sec: 18.7077
INFO:tensorflow:step = 801, loss = 0.08796086 (5.345 sec)
INFO:tensorflow:global_step/sec: 19.2725
INFO:tensorflow:step = 901, loss = 0.078957975 (5.189 sec)
INFO:tensorflow:Saving checkpoints for 1000 into outputdir/model.ckp
t.
INFO:tensorflow:Loss for final step: 0.07356742.
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Starting evaluation at 2018-09-12-20:03:03
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from outputdir/model.ckpt-1000
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Finished evaluation at 2018-09-12-20:03:04

```

```
INFO:tensorflow:Saving dict for global step 1000: global_step = 1000, loss = 0.06301524, rmse = 0.25102836
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Signatures INCLUDED in export for Classify: None
INFO:tensorflow:Signatures INCLUDED in export for Regress: None
INFO:tensorflow:Signatures INCLUDED in export for Predict: ['predict_export_outputs', 'serving_default']
INFO:tensorflow:Restoring parameters from outputdir/model.ckpt-1000
INFO:tensorflow:Assets added to graph.
INFO:tensorflow:No assets to write.
INFO:tensorflow:SavedModel written to: b"outputdir/export/exporter/temp-b'1536782584'/saved_model.pb"
```

Try out online prediction. This is how the REST API will work after you train on Cloud ML Engine

In [39]:

```
%%writefile test.json
{"rawdata_input": [0,0.214,0.406,0.558,0.655,0.687,0.65,0.549,0.393]}
```

Overwriting test.json

In [40]:

```
# local predict doesn't work with Python 3 yet.
# %%bash
# MODEL_DIR=$(ls ./outputdir/export/exporter/)
# gcloud ml-engine local predict --model-dir=./outputdir/export/exporter/$MODEL_DIR --json-instances=test.json
```

Cloud ML Engine

Now to train on Cloud ML Engine.

In [41]:

```
%%bash
# Run module on Cloud ML Engine
OUTDIR=gs://${BUCKET}/simplernn/model_trained
JOBNAME=simplernn_$(date -u +%Y%m%d_%H%M%S)
gsutil -m rm -rf $OUTDIR
gcloud ml-engine jobs submit training $JOBNAME \
  --region=$REGION \
  --module-name=trainer.task \
  --package-path=${PWD}/simplernn/trainer \
  --job-dir=$OUTDIR \
  --staging-bucket=gs://${BUCKET} \
  --scale-tier=BASIC \
  --runtime-version=1.4 \
  -- \
  --train_data_paths="gs://${BUCKET}/train.csv*" \
  --eval_data_paths="gs://${BUCKET}/valid.csv*" \
  --output_dir=$OUTDIR
```

jobId: simplernn_180912_200305

state: QUEUED

CommandException: 1 files/objects could not be removed.

Job [simplernn_180912_200305] submitted successfully.

Your job is still active. You may view the status of your job with the command

```
$ gcloud ml-engine jobs describe simplernn_180912_200305
```

or continue streaming the logs with the command

```
$ gcloud ml-engine jobs stream-logs simplernn_180912_200305
```

Variant: long sequence

To create short sequences from a very long sequence.

In [42]:

```
import tensorflow as tf
import numpy as np

def breakup(sess, x, lookback_len):
    N = sess.run(tf.size(x))
    windows = [tf.slice(x, [b], [lookback_len]) for b in range(0, N-lookback_len)]
    windows = tf.stack(windows)
    return windows

x = tf.constant(np.arange(1,11, dtype=np.float32))
with tf.Session() as sess:
    print('input=', x.eval())
    seqx = breakup(sess, x, 5)
    print('output=', seqx.eval())
```

```
input= [ 1.  2.  3.  4.  5.  6.  7.  8.  9. 10.]
output= [[1. 2. 3. 4. 5.]
 [2. 3. 4. 5. 6.]
 [3. 4. 5. 6. 7.]
 [4. 5. 6. 7. 8.]
 [5. 6. 7. 8. 9.]]
```

Variant: Keras

You can also invoke a Keras model from within the Estimator framework by creating an estimator from the compiled Keras model:

In [43]:

```
def make_keras_estimator(output_dir):
    from tensorflow import keras
    model = keras.models.Sequential()
    model.add(keras.layers.Dense(32, input_shape=(N_INPUTS,), name=TIMESERIES_INPU
T_LAYER))
    model.add(keras.layers.Activation('relu'))
    model.add(keras.layers.Dense(1))
    model.compile(loss = 'mean_squared_error',
                  optimizer = 'adam',
                  metrics = ['mae', 'mape']) # mean absolute [percentage] error
    return keras.estimator.model_to_estimator(model)
```

In []:

```
%%bash
# Run module as-is
echo $PWD
rm -rf outputdir
export PYTHONPATH=${PYTHONPATH}:${PWD}/simplernn
python -m trainer.task \
  --train_data_paths="${PWD}/train.csv*" \
  --eval_data_paths="${PWD}/valid.csv*" \
  --output_dir=${PWD}/outputdir \
  --job-dir=./tmp --keras
```


Copyright 2017 Google Inc. Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at <http://www.apache.org/licenses/LICENSE-2.0> (<http://www.apache.org/licenses/LICENSE-2.0>) Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License